



Lab02

Crowd Counting

PyTorch tutorial

- Official tutorial
 - <https://pytorch.org/tutorials/>
- 莫凡
 - <https://mofanpy.com/tutorials/machine-learning/torch/>
- AssemblyAI - PyTorch Crash Course
 - <https://www.youtube.com/watch?v=OlenNRt2bjg>

You can only use PyTorch in this Lab!!

Crowd Counting

- **Crowd counting** is a computer vision technique that aims to estimate the number of people in crowded images using deep learning models.
- It is essential for analyzing large gatherings where manual counting is impractical.
- Applications: Public safety monitoring, event management, smart city planning, transportation hubs (metro, airports), retail analytics, and disaster response.

Dataset

- UCSD Pedestrian Dataset
- Image size: 238*158 grayscale
- Ground Truth: 3 values for counts of walking away, toward, and total
- Training: 2500 Validation: 700 Testing: 800 (200 public + 600 private)




Task 1 of This Lab

- In “Lab02_CDenseNet.ipynb”
 - Build CDenseNet by yourself
 - Achieve MAE of **2.4 or lower** on public testing data
(Put the screenshot in your report)

CDenseNet


- DenseNet concatenates all prior features \rightarrow channels explode \rightarrow heavy compute/weights.
- Lightweight Dense Block (LDB)** fuse by element-wise sum; use concat only at block input/output \rightarrow fixed channel width with reuse.



(d) Lightweight dense block

CDenseNet

- Simplified Compress DenseNet (CDenseNet)
- Choose $t = 0.5$ & $n = 16$



Please follow this model architecture!
We will check your implementation

Task 1 of This Lab

- Finish these parts (CDenseNet.py & training flow).

```

1 import torch
2 import torch.nn as nn
3
4 class LDB(nn.Module):
5     def __init__(self, in_channel: int, t: float = 0.5):
6         pass
7
8     def forward(self, x):
9         pass
10
11 class CDenseNet(nn.Module):
12     def __init__(self, n: int = 16, t: float = 0.5):
13         pass
14
15     def forward(self, x):
16         pass

```

```

4 ##### implement your optimizer #####
5 ## you can use any training methods if you want (ex:lr decay, weight decay....)
6
7 # you can try 10~15 at first
8 num_epochs =
9 # Learning rate
10 lr =
11
12 # Loss function
13 criterion =
14 # Optimizer
15 optimizer =
16 # Learning rate scheduler (optional)
17 scheduler = None

```

```

9 for epoch in range(1, num_epochs + 1):
10     # ----- Training phase -----
11     model.train() # Set the model to training mode
12     running_loss = 0.0
13     train_bar = tqdm(train_loader, desc=f'Epoch {epoch}/{num_epochs} [Train]', leave=False, position=0, smoothing=0.1)
14     for in_img, people_cnts in train_bar:
15         in_img, people_cnts = in_img.to(device, non_blocking=True), people_cnts.to(device, non_blocking=True)
16
17         #####
18         # Please finish the "Training phase" code here.
19
20
21         #####
22
23         running_loss += loss.item() * people_cnts.size(0)
24         train_bar.set_postfix(loss=f'{loss.item():.4f}')
25
26     # ----- Validation phase -----
27     model.eval() # Set the model to evaluation mode
28     val_loss = 0
29     # Per-component MAE/RMSE accumulators for [r, l, t]
30     abs_sum = torch.zeros(3, dtype=torch.float64)
31     sqr_sum = torch.zeros(3, dtype=torch.float64)
32
33     with torch.no_grad():
34         val_bar = tqdm(val_loader, desc=f'Epoch {epoch}/{num_epochs} [Val]', leave=False, position=0, smoothing=0.1)
35         for in_img, people_cnts in val_bar:
36             in_img, people_cnts = in_img.to(device, non_blocking=True), people_cnts.to(device, non_blocking=True)
37
38             #####
39             # Forward pass for validation
40             # Please finish the "Validation phase" code here.
41
42
43             #####
44
45             # Calculate metrics for validation results
46             err = outputs - people_cnts
47             abs_sum += err.abs().sum(dim=0).double().cpu()
48             sqr_sum += (err ** 2).sum(dim=0).double().cpu()
49             val_bar.set_postfix(loss=f'{loss.item():.4f}')

```


Task 2 of This Lab

- In Task2
 - Do your best to improve the prediction accuracy
 - Calling different models with pretrained weight is allowed
 - Basically, any methods you learn are allowed
 - Achieve MAE of **2.0 or lower** on public testing data
(put the screenshot in your report)

Report

- Your report should include/answer
 - Required
 - Screenshot of Task 1 (MAE on public testing data ≤ 2.4)
 - Screenshot of Task 2 (MAE on public testing data ≤ 2.0)
 - In Task 2
 - What model did you choose?
 - Why did you choose this model? What advantages does it offer?
 - Compare the characteristics of **MAE (Mean Absolute Error)** and **RMSE (Root Mean Square Error)**. In what types of scenarios might one be preferred over the other?
 - Another popular method in crowd counting is “**density map estimation**.” Briefly explain what density map estimation means in the context of crowd counting. How does it differ from the regression-based approach used in our implementation? Give at least one metrics used to evaluate it.
 - Can include but not limited to
 - Anything you do to improve the accuracy.
 - Discuss any challenges you faced.

Score

- MAE on public testing data in Task 1 ≤ 2.4 (30%)
 - If the model architecture in Task 1 is incorrect, points will be deducted accordingly
- MAE on public testing data in Task 2 ≤ 2.0 (30%)
- Report (30%)
- Performance ranking for Task 1 (10%)
 - Ranked based on MAE on the full testing data in Task 1
 - 0 points will be given if your model is found trained in an abnormal way
- Please do not plagiarize, or you will receive 0 points if caught

Reminder

- Submit Deadline : 2 week (2025-10-06 23:59)
- Upload these files to E3
 - Lab02_CDenseNet_StudentID.ipynb
 - CDenseNet_StudentID.py
 - model_StudentID.pth (of Task 1)
 - summary_StudentID.txt (of Task 1)
 - Lab02_report_StudentID.pdf

Supplements

- **paper**
 - <https://arxiv.org/abs/1912.07016>
 - <https://ieeexplore.ieee.org/document/6054049>

HAVE FUN !!!